

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv(r"E:\154\C6_bmi - C6_bmi.csv").dropna()
df
```

Out[3]:

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...	...	...	...	...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

500 rows × 4 columns

```
In [4]: df.head()
```

Out[4]:

	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3

In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 500 entries, 0 to 499
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Gender  500 non-null      object
1   Height  500 non-null      int64
2   Weight  500 non-null      int64
3   Index   500 non-null      int64
dtypes: int64(3), object(1)
memory usage: 19.5+ KB
```

In [6]: df.describe()

Out[6]:

	Height	Weight	Index
<b>count</b>	500.000000	500.000000	500.000000
<b>mean</b>	169.944000	106.000000	3.748000
<b>std</b>	16.375261	32.382607	1.355053
<b>min</b>	140.000000	50.000000	0.000000
<b>25%</b>	156.000000	80.000000	3.000000
<b>50%</b>	170.500000	106.000000	4.000000
<b>75%</b>	184.000000	136.000000	5.000000
<b>max</b>	199.000000	160.000000	5.000000

In [7]: df.columns

Out[7]: Index(['Gender', 'Height', 'Weight', 'Index'], dtype='object')

In [8]: feature\_matrix = df[['Height', 'Weight']]  
target\_vector = df[['Index']]

In [9]: fs=StandardScaler().fit\_transform(feature\_matrix)  
logr=LogisticRegression()  
logr.fit(fs,target\_vector)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().  
return f(\*args, \*\*kwargs)

Out[9]: LogisticRegression()

In [10]: observation=[[1,2]]

```
In [11]: prediction=logr.predict(observation)
print(prediction)

[5]
```

```
In [12]: logr.classes_
```

```
Out[12]: array([0, 1, 2, 3, 4, 5], dtype=int64)
```

```
In [13]: logr.predict_proba(observation)[0][0]
```

```
Out[13]: 5.5956697582538237e-11
```

```
In [14]: logr.predict_proba(observation)[0][1]
```

```
Out[14]: 6.059900360819463e-10
```

## Random Forest

```
In [37]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [38]: df=pd.read_csv(r"E:\154\C6_bmi - C6_bmi.csv")
```

```
In [39]: df['Gender'].value_counts()
```

```
Out[39]: Female    255
Male          245
Name: Gender, dtype: int64
```

```
In [40]: x=df[['Height','Weight']]
y=df['Gender']
```

```
In [41]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [42]: from sklearn.ensemble import RandomForestClassifier
rfc =RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[42]: RandomForestClassifier()
```

```
In [43]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]  
                    }
```

```
In [44]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac  
grid_search.fit(x_train,y_train)
```

```
Out[44]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                    param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                'min_samples_leaf': [5, 10, 15, 20, 25],  
                                'n_estimators': [10, 20, 30, 40, 50]},  
                    scoring='accuracy')
```

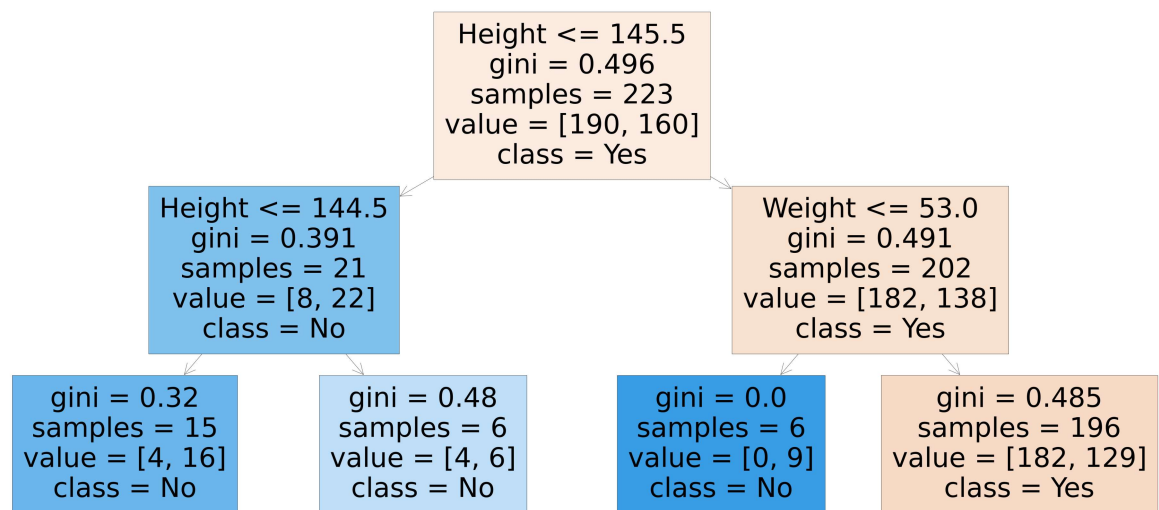
```
In [45]: grid_search.best_score_
```

```
Out[45]: 0.5314285714285714
```

```
In [46]: rfc_best=grid_search.best_estimator_
```

```
In [47]: from sklearn.tree import plot_tree
plt.figure(figsize=(89,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns, class_names=["Yes",
    )
```

```
Out[47]: [Text(2483.1, 1812.0, 'Height <= 145.5\ngini = 0.496\nsamples = 223\nvalue =
[190, 160]\n\nclass = Yes'),
Text(1241.55, 1087.2, 'Height <= 144.5\ngini = 0.391\nsamples = 21\nvalue =
[8, 22]\n\nclass = No'),
Text(620.775, 362.39999999999986, 'gini = 0.32\nsamples = 15\nvalue = [4, 1
6]\n\nclass = No'),
Text(1862.3249999999998, 362.39999999999986, 'gini = 0.48\nsamples = 6\nvalu
e = [4, 6]\n\nclass = No'),
Text(3724.6499999999996, 1087.2, 'Weight <= 53.0\ngini = 0.491\nsamples = 20
2\nvalue = [182, 138]\n\nclass = Yes'),
Text(3103.875, 362.39999999999986, 'gini = 0.0\nsamples = 6\nvalue = [0, 9]
\n\nclass = No'),
Text(4345.425, 362.39999999999986, 'gini = 0.485\nsamples = 196\nvalue = [18
2, 129]\n\nclass = Yes')]
```



In [ ]: