

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df=pd.read_csv(r"E:\154\C4_framingham - C4_framingham.csv").dropna()
df
```

Out[3]:

	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	t
1	39	4.0	0	0.0	0.0	0	0	0	
0	46	2.0	0	0.0	0.0	0	0	0	
1	48	1.0	1	20.0	0.0	0	0	0	
0	61	3.0	1	30.0	0.0	0	1	0	
0	46	3.0	1	23.0	0.0	0	0	0	
..	...	...	...	...	...	...	...	...	
1	58	3.0	0	0.0	0.0	0	1	0	
1	68	1.0	0	0.0	0.0	0	1	0	
1	50	1.0	1	1.0	0.0	0	1	0	
1	51	3.0	1	43.0	0.0	0	0	0	
0	52	2.0	0	0.0	0.0	0	0	0	

× 16 columns



```
In [4]: df.head()
```

Out[4]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes
0	1	39	4.0	0	0.0	0.0	0	0	
1	0	46	2.0	0	0.0	0.0	0	0	
2	1	48	1.0	1	20.0	0.0	0	0	
3	0	61	3.0	1	30.0	0.0	0	1	
4	0	46	3.0	1	23.0	0.0	0	0	



In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3656 entries, 0 to 4237
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   male                  3656 non-null   int64  
 1   age                   3656 non-null   int64  
 2   education             3656 non-null   float64 
 3   currentSmoker         3656 non-null   int64  
 4   cigsPerDay            3656 non-null   float64 
 5   BPMeds                3656 non-null   float64 
 6   prevalentStroke       3656 non-null   int64  
 7   prevalentHyp          3656 non-null   int64  
 8   diabetes              3656 non-null   int64  
 9   totChol              3656 non-null   float64 
10   sysBP                3656 non-null   float64 
11   diaBP                3656 non-null   float64 
12   BMI                  3656 non-null   float64 
13   heartRate            3656 non-null   float64 
14   glucose              3656 non-null   float64 
15   TenYearCHD           3656 non-null   int64  
dtypes: float64(9), int64(7)
memory usage: 485.6 KB
```

In [6]: `df.describe()`

Out[6]:

cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	
356.000000	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000	3656.000000
9.022155	0.030361	0.005744	0.311543	0.027079	236.873085	132.368025	8.000000
11.918869	0.171602	0.075581	0.463187	0.162335	44.096223	22.092444	1.000000
0.000000	0.000000	0.000000	0.000000	0.000000	113.000000	83.500000	4.000000
0.000000	0.000000	0.000000	0.000000	0.000000	206.000000	117.000000	7.000000
0.000000	0.000000	0.000000	0.000000	0.000000	234.000000	128.000000	8.000000
20.000000	0.000000	0.000000	1.000000	0.000000	263.250000	144.000000	9.000000
70.000000	1.000000	1.000000	1.000000	1.000000	600.000000	295.000000	14.000000

In [7]: `df.columns`

Out[7]: Index(['male', 'age', 'education', 'currentSmoker', 'cigsPerDay', 'BPMeds', 'prevalentStroke', 'prevalentHyp', 'diabetes', 'totChol', 'sysBP', 'diaBP', 'BMI', 'heartRate', 'glucose', 'TenYearCHD'], dtype='object')

```
In [8]: feature_matrix = df.iloc[:,0:15]
target_vector = df.iloc[:, -1]
```

```
In [9]: fs=StandardScaler().fit_transform(feature_matrix)
logr=LogisticRegression()
logr.fit(fs,target_vector)
```

Out[9]: LogisticRegression()

```
In [10]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15]]
```

```
In [11]: prediction=logr.predict(observation)
print(prediction)
```

[1]

```
In [12]: logr.classes_
```

Out[12]: array([0, 1], dtype=int64)

```
In [13]: logr.predict_proba(observation)[0][0]
```

Out[13]: 0.0002214783507201723

```
In [14]: logr.predict_proba(observation)[0][1]
```

Out[14]: 0.9997785216492798

## Random Forest

```
In [15]: df['TenYearCHD'].value_counts()
```

Out[15]: 0 3099  
1 557  
Name: TenYearCHD, dtype: int64

```
In [16]: x=df[['TenYearCHD']]
y=df['TenYearCHD']
```

```
In [18]: g1={'TenYearCHD':{'0':1, "1":2}}
df=df.replace(g1)
df
```

Out[18]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp
0	1	39	4.0	0	0.0	0.0	0	0
1	0	46	2.0	0	0.0	0.0	0	0
2	1	48	1.0	1	20.0	0.0	0	0
3	0	61	3.0	1	30.0	0.0	0	1
4	0	46	3.0	1	23.0	0.0	0	0
...	...	...	...	...	...	...	...	...
4231	1	58	3.0	0	0.0	0.0	0	1
4232	1	68	1.0	0	0.0	0.0	0	1
4233	1	50	1.0	1	1.0	0.0	0	1
4234	1	51	3.0	1	43.0	0.0	0	0
4237	0	52	2.0	0	0.0	0.0	0	0

3656 rows × 16 columns

```
In [24]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [25]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[25]: RandomForestClassifier()

```
In [26]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                        'n_estimators': [10,20,30,40,50]
                        }
```

```
In [27]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="a
grid_search.fit(x_train,y_train)
```

```
Out[27]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

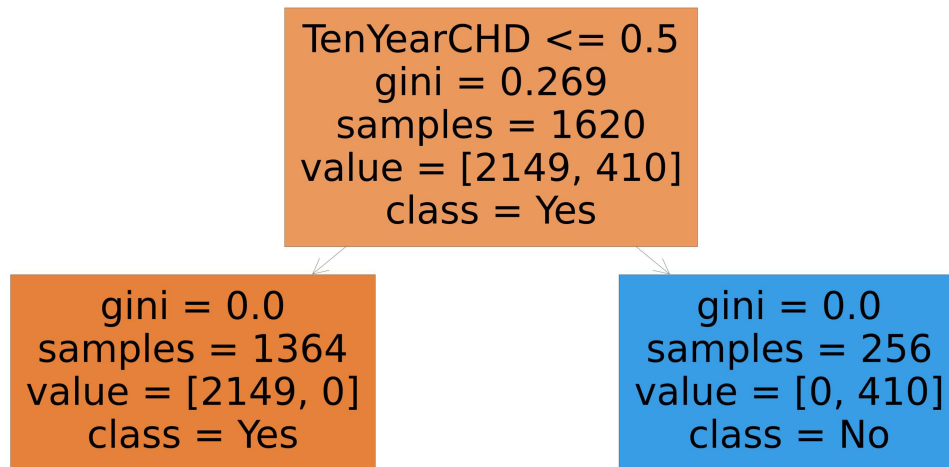
```
In [28]: grid_search.best_score_
```

Out[28]: 1.0

```
In [29]: rfc_best = grid_search.best_estimator_
```

```
In [31]: from sklearn.tree import plot_tree  
plt.figure(figsize=(95,40))  
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['Yes'
```

```
Out[31]: [Text(2650.5, 1630.8000000000002, 'TenYearCHD <= 0.5\ngini = 0.269\nsamples =  
1620\nvalue = [2149, 410]\nclass = Yes'),  
Text(1325.25, 543.5999999999999, 'gini = 0.0\nsamples = 1364\nvalue = [2149,  
0]\nclass = Yes'),  
Text(3975.75, 543.5999999999999, 'gini = 0.0\nsamples = 256\nvalue = [0, 41  
0]\nclass = No')]
```



```
In [ ]:
```