

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [ ]:
```

```
In [2]: df=pd.read_csv(r"E:\154\1_ionosphere - 1_ionosphere.csv")
df
```

```
Out[2]:
```

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.0376	...	-0.511
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.265
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.402
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.906
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.651
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.015
...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-0.042
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	0.013
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	0.031
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-0.020
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-0.157

350 rows × 35 columns



```
In [3]: df.head()
```

```
Out[3]:
```

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.0376	...	-0.51171
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.26569
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.40220
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.90695
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.65158
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.01535

5 rows × 35 columns



Data Cleaning and Data Preprocessing

In [4]: df.info()


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 350 entries, 0 to 349
Data columns (total 35 columns):
#   Column          Non-Null Count  Dtype
---  -
0   1                350 non-null    int64
1   0                350 non-null    int64
2   0.99539          350 non-null    float64
3   -0.05889         350 non-null    float64
4   0.85243          350 non-null    float64
5   0.02306          350 non-null    float64
6   0.83398          350 non-null    float64
7   -0.37708         350 non-null    float64
8   1.1              350 non-null    float64
9   0.0376           350 non-null    float64
10  0.85243.1        350 non-null    float64
11  -0.17755         350 non-null    float64
12  0.59755          350 non-null    float64
13  -0.44945         350 non-null    float64
14  0.60536          350 non-null    float64
15  -0.38223         350 non-null    float64
16  0.84356          350 non-null    float64
17  -0.38542         350 non-null    float64
18  0.58212          350 non-null    float64
19  -0.32192         350 non-null    float64
20  0.56971          350 non-null    float64
21  -0.29674         350 non-null    float64
22  0.36946          350 non-null    float64
23  -0.47357         350 non-null    float64
24  0.56811          350 non-null    float64
25  -0.51171         350 non-null    float64
26  0.41078          350 non-null    float64
27  -0.46168         350 non-null    float64
28  0.21266          350 non-null    float64
29  -0.3409          350 non-null    float64
30  0.42267          350 non-null    float64
31  -0.54487         350 non-null    float64
32  0.18641          350 non-null    float64
33  -0.453           350 non-null    float64
34  g                350 non-null    object
dtypes: float64(32), int64(2), object(1)
memory usage: 95.8+ KB
```

In [5]: `df.describe()`

Out[5]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.3770
count	350.000000	350.0	350.000000	350.000000	350.000000	350.000000	350.000000	350.000000
mean	0.891429	0.0	0.640330	0.044667	0.600350	0.116154	0.549284	0.12077
std	0.311546	0.0	0.498059	0.442032	0.520431	0.461443	0.493124	0.52081
min	0.000000	0.0	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000
25%	1.000000	0.0	0.471517	-0.065388	0.412555	-0.024868	0.209105	-0.05348
50%	1.000000	0.0	0.870795	0.016700	0.808620	0.021170	0.728000	0.01508
75%	1.000000	0.0	1.000000	0.194727	1.000000	0.335317	0.970445	0.45157
max	1.000000	0.0	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 34 columns



In [6]: `df.columns`

Out[6]: Index(['1', '0', '0.99539', '-0.05889', '0.85243', '0.02306', '0.83398',
'-0.37708', '1.1', '0.0376', '0.85243.1', '-0.17755', '0.59755',
'-0.44945', '0.60536', '-0.38223', '0.84356', '-0.38542', '0.58212',
'-0.32192', '0.56971', '-0.29674', '0.36946', '-0.47357', '0.56811',
'-0.51171', '0.41078', '-0.46168', '0.21266', '-0.3409', '0.42267',
'-0.54487', '0.18641', '-0.453', 'g'],
dtype='object')

In [7]: `feature_matrix = df.iloc[:,0:34]`
`target_vector = df.iloc[:, -1]`

In [8]: `fs = StandardScaler().fit_transform(feature_matrix)`
`logr = LogisticRegression()`
`logr.fit(fs, target_vector)`

Out[8]: `LogisticRegression()`

```
In [9]: observation=[[1.0,0.0,1.0,-0.18829,0.93035,  
-0.36156,  
-0.10868,  
-0.93597,  
1.0,  
-0.04549,  
0.50874,  
-0.67743,  
0.34432,  
-0.69707,  
-0.51685,  
-0.97515,  
0.05499,  
-0.62237,  
0.33109,  
-1.0,  
-0.13151,  
-0.453,  
-0.18056,  
-0.35734,  
-0.20332,  
-0.26569,  
-0.20468,  
-0.18401,  
-0.1904,  
-0.11593,  
-0.16626,  
-0.06288,  
-0.13738,  
-0.02447]]  
prediction = logr.predict(observation)  
print(prediction)
```

```
['g']
```

```
In [10]: logr.classes_
```

```
Out[10]: array(['b', 'g'], dtype=object)
```

```
In [11]: logr.predict_proba(observation)
```

```
Out[11]: array([[0.07006552, 0.92993448]])
```

Random Forest

```
In [12]: df['g'].value_counts()
```

```
Out[12]: g    224  
b    126  
Name: g, dtype: int64
```

```
In [13]: x=df.drop('g', axis=1)
y=df['g']
```

```
In [14]: g1={"g":{"g":1, "b":2}}
df=df.replace(g1)
df
```

Out[14]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.0376	...	-0.511
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-0.265
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-0.402
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	0.906
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-0.651
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-0.015
...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-0.042
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	0.013
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	0.031
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-0.020
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-0.157

350 rows × 35 columns



```
In [15]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [16]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[16]: RandomForestClassifier()

```
In [17]: parameters = {'max_depth':[1,2,3,4,5], 'min_samples_leaf':[5,10,15,20,25],
                        'n_estimators': [10,20,30,40,50]}
}
```

```
In [18]: from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[18]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')

```
In [19]: grid_search.best_score_
```

```
Out[19]: 0.9385245901639344
```

```
In [20]: rfc_best = grid_search.best_estimator_
```

```
In [21]: from sklearn.tree import plot_tree
plt.figure(figsize=(89,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['Yes',
```

```
Out[21]: [Text(2031.6272727272726, 1956.96, '-0.44945 <= -0.549\ngini = 0.441\nsamples = 162\nvalue = [80, 164]\nnclass = No'),
Text(902.9454545454545, 1522.0800000000002, '-0.05889 <= -0.253\ngini = 0.198\nsamples = 13\nvalue = [24, 3]\nnclass = Yes'),
Text(451.47272727272724, 1087.2, 'gini = 0.444\nsamples = 5\nvalue = [6, 3]\nnclass = Yes'),
Text(1354.4181818181817, 1087.2, 'gini = 0.0\nsamples = 8\nvalue = [18, 0]\nnclass = Yes'),
Text(3160.3090909090906, 1522.0800000000002, '-0.46168 <= -0.926\ngini = 0.383\nsamples = 149\nvalue = [56, 161]\nnclass = No'),
Text(2257.3636363636363, 1087.2, '-0.47357 <= -0.093\ngini = 0.219\nsamples = 13\nvalue = [14, 2]\nnclass = Yes'),
Text(1805.890909090909, 652.3200000000002, 'gini = 0.408\nsamples = 5\nvalue = [5, 2]\nnclass = Yes'),
Text(2708.8363636363633, 652.3200000000002, 'gini = 0.0\nsamples = 8\nvalue = [9, 0]\nnclass = Yes'),
Text(4063.2545454545454, 1087.2, '-0.17755 <= 0.984\ngini = 0.331\nsamples = 136\nvalue = [42, 159]\nnclass = No'),
Text(3611.7818181818181, 652.3200000000002, '-0.05889 <= -0.457\ngini = 0.293\nsamples = 130\nvalue = [34, 157]\nnclass = No'),
Text(3160.3090909090906, 217.44000000000005, 'gini = 0.0\nsamples = 5\nvalue = [5, 0]\nnclass = Yes'),
Text(4063.2545454545454, 217.44000000000005, 'gini = 0.263\nsamples = 125\nvalue = [29, 157]\nnclass = No'),
Text(4514.727272727272, 652.3200000000002, 'gini = 0.32\nsamples = 6\nvalue = [8, 2]\nnclass = Yes')]
```

