

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: df_train=pd.read_csv(r"E:\154\C8_loan-train - C8_loan-train.csv")
df_test=pd.read_csv(r"E:\154\C8_loan-test - C8_loan-test.csv")
df_train
```

Out[3]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coap
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	
...	...	...	...	...	...	...	...	
609	LP002978	Female	No	0	Graduate	No	2900	
610	LP002979	Male	Yes	3+	Graduate	No	4106	
611	LP002983	Male	Yes	1	Graduate	No	8072	
612	LP002984	Male	Yes	2	Graduate	No	7583	
613	LP002990	Female	No	0	Graduate	Yes	4583	

614 rows × 13 columns



In [4]: df\_test

Out[4]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001015	Male	Yes	0	Graduate	No	5720	
1	LP001022	Male	Yes	1	Graduate	No	3076	
2	LP001031	Male	Yes	2	Graduate	No	5000	
3	LP001035	Male	Yes	2	Graduate	No	2340	
4	LP001051	Male	No	0	Not Graduate	No	3276	
...	...	...	...	...	...	...	...	...
362	LP002971	Male	Yes	3+	Not Graduate	Yes	4009	
363	LP002975	Male	Yes	0	Graduate	No	4158	
364	LP002980	Male	No	0	Graduate	No	3250	
365	LP002986	Male	Yes	0	Graduate	No	5000	

In [5]: df1=df\_train.dropna()

In [6]: df2=df\_test.dropna()

In [7]: df1.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 480 entries, 1 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                480 non-null    object
1   Gender                 480 non-null    object
2   Married                480 non-null    object
3   Dependents             480 non-null    object
4   Education              480 non-null    object
5   Self_Employed          480 non-null    object
6   ApplicantIncome        480 non-null    int64
7   CoapplicantIncome      480 non-null    float64
8   LoanAmount             480 non-null    float64
9   Loan_Amount_Term       480 non-null    float64
10  Credit_History         480 non-null    float64
11  Property_Area          480 non-null    object
12  Loan_Status            480 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 52.5+ KB
```

In [8]: df1.columns

Out[8]: Index(['Loan\_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self\_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan\_Amount\_Term', 'Credit\_History', 'Property\_Area', 'Loan\_Status'], dtype='object')

```
In [9]: feature_matrix=df1[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
    'Loan_Amount_Term', 'Credit_History']]
target_vector=df1[['Self_Employed']]
```

```
In [10]: feature_matrix.shape
```

```
Out[10]: (480, 5)
```

```
In [11]: target_vector.shape
```

```
Out[11]: (480, 1)
```

```
In [12]: from sklearn.preprocessing import StandardScaler
```

```
In [13]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [14]: logr=LogisticRegression()
logr.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
return f(*args, **kwargs)
```

```
Out[14]: LogisticRegression()
```

```
In [15]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 289 entries, 0 to 366
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                289 non-null    object
1   Gender                 289 non-null    object
2   Married                289 non-null    object
3   Dependents             289 non-null    object
4   Education              289 non-null    object
5   Self_Employed          289 non-null    object
6   ApplicantIncome        289 non-null    int64
7   CoapplicantIncome      289 non-null    int64
8   LoanAmount             289 non-null    float64
9   Loan_Amount_Term       289 non-null    float64
10  Credit_History         289 non-null    float64
11  Property_Area          289 non-null    object
dtypes: float64(3), int64(2), object(7)
memory usage: 29.4+ KB
```

```
In [16]: df2.columns
```

```
Out[16]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
    'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
    'Loan_Amount_Term', 'Credit_History', 'Property_Area'],
    dtype='object')
```

```
In [17]: observation=df2[['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
                        'Loan_Amount_Term', 'Credit_History']]
```

```
In [18]: prediction=logr.predict(observation)
print(prediction)
```

```
['Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes' 'Yes'
'Yes']
```

```
In [19]: logr.classes_
```

```
Out[19]: array(['No', 'Yes'], dtype=object)
```



```
In [24]: g1={'Self_Employed':{'No':1, 'Yes':2}}
df2=df2.replace(g1)
df2
```

Out[24]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coap
0	LP001015	Male	Yes	0	Graduate	1	5720	
1	LP001022	Male	Yes	1	Graduate	1	3076	
2	LP001031	Male	Yes	2	Graduate	1	5000	
4	LP001051	Male	No	0	Not Graduate	1	3276	
5	LP001054	Male	Yes	0	Not Graduate	2	2165	
...	...	...	...	...	...	...	...	...
361	LP002969	Male	Yes	1	Graduate	1	2269	
362	LP002971	Male	Yes	3+	Not Graduate	2	4009	
363	LP002975	Male	Yes	0	Graduate	1	4158	
365	LP002986	Male	Yes	0	Graduate	1	5000	
366	LP002989	Male	No	0	Graduate	2	9200	

289 rows × 12 columns



```
In [25]: from sklearn.model_selection import train_test_split
```

```
In [26]: x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.70)
```

```
In [27]: from sklearn.ensemble import RandomForestClassifier
```

```
In [28]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[28]: RandomForestClassifier()

```
In [29]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
}
```

```
In [30]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac
grid_search.fit(x_train,y_train)
```

Out[30]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 3, 4, 5],  
'min\_samples\_leaf': [5, 10, 15, 20, 25],  
'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

```
In [31]: grid_search.best_score_
```

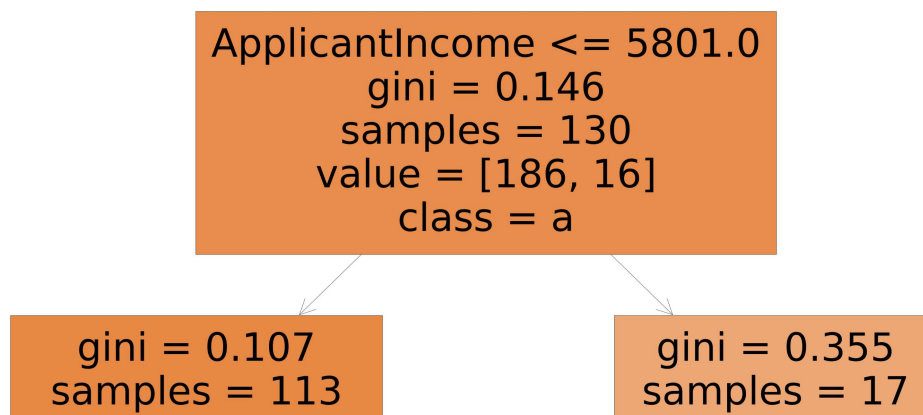
```
Out[31]: 0.900990099009901
```

```
In [32]: rfc_best=grid_search.best_estimator_
```

```
In [33]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))  
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b'])
```

```
Out[33]: [Text(2232.0, 1630.8000000000002, 'ApplicantIncome <= 5801.0\ngini = 0.146  
\nsamples = 130\nvalue = [186, 16]\nnclass = a'),  
Text(1116.0, 543.5999999999999, 'gini = 0.107\nsamples = 113\nvalue = [16  
6, 10]\nnclass = a'),  
Text(3348.0, 543.5999999999999, 'gini = 0.355\nsamples = 17\nvalue = [20,  
6]\nnclass = a')]
```



```
In [ ]:
```